

UNIVERSITAT POLITECNICA DE
CATALUNYA

PARALLELISM

*Lab 4: Divide and Conquer parallelism with
OpenMP: Sorting*

*Roger Vilaseca Darne and Xavier Martin Ballesteros
PAR4110*



15th May 2019, Q2

Contents

1	Introduction	2
2	Analysis with <i>Tareador</i>	3
2.1	Merge Function	3
2.2	Multisort Function	3
2.3	Task dependency graph (TDG)	4
2.4	Parallel Performance and Scalability	5
3	Parallelization and performance analysis with tasks	6
3.1	Leaf Strategy	7
3.2	Tree Strategy	9
3.3	Task Cut-Off Mechanism	13
4	Parallelization and performance analysis with dependent tasks	17
5	Optionals	19
5.1	Optional 1	19
5.1.1	Boada-1 to 4	20
5.1.2	Boada-5	20
5.1.3	Boada-6 to 8	21
5.2	Optional 2	22
6	Conclusion	24
7	Annex	25
7.1	Scalability analysis: Traces using different number of processors	25

1 Introduction

Sorting is such an important part in Algorithm that it is crucial to find the best algorithms that sort vectors in the least possible execution time.

Mergesort is a sorting algorithm which combines a "divide and conquer" strategy. It divides the initial vector into multiple subvectors recursively and when the size of the subvectors reaches a value, a sequential quicksort is applied to each of the created subvectors. Finally, with a merge of the subvectors we back into a single sorted vector.

In this three sessions of the laboratory we will analyse the performance of this algorithm using different parallel strategies. The first one will be the Leaf strategy, in which we create a task each time we are in the base case of the function. The second one is the Tree strategy, in which we define a task during the recursive decomposition.

To do this analysis we will first use the *Tareador* tool to see whether we can use both strategies or not, depending on the dependencies that exist in the program tasks. Once we have done this, we will parallelise the code using the *OpenMP* clauses.

First we will implement the Leaf strategy and later the Tree one. For the Tree strategy we will use different versions of its implementation (with and without cut-off and using the *depend* clause).

To be able to analyse the parallelisation and performance of the program we will be using the *Paraver* tool and the scripts that we have been given.

2 Anlysis with *Tareador*

In order to study which is the best approach that we should use to parallelize the code (using the *OpenMP* directives), a study about the task dependencies is needed.

In the following subsections, we will analyse the functions *merge* and *multisort* of the given code using *Tareador* to see the potential parallelism they have and also their dependences. As both functions are recursive, it is very likely that we will be able to parallelise them.

The modified code can be found in the *multisort-tareador.c* file, inside the codes directory.

2.1 Merge Function

The merge function is responsible for sorting two vectors, merging them into one. This is done using recursive calls that divide the vectors into smaller ones.

Our study in this function focuses on the two recursive calls done in the else fragment of the function. A task is created every time we make a recursive call to the function. As a consequence, some created tasks will also execute the basicmerge function.

To do that, we have used the *tareador_start_task* and *tareador_end_task* functions. The resulting code is shown below.

```
void merge(long n, T left[n], T right[n], T result[n*2], long start,
           long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        tareador_start_task("mef1");
        merge(n, left, right, result, start, length/2);
        tareador_end_task("mef1");

        tareador_start_task("mef2");
        merge(n, left, right, result, start + length/2, length/2);
        tareador_end_task("mef2");
    }
}
```

Figure 1: Modified version of the merge function.

2.2 Multisort Function

The multisort function divides the input vector into 4 smaller vectors, apply a recursive call for each of them and then merges the 4 vectors into a big sorted one.

We will use the same strategy than in the merge function. This is, creating a task for each recursive call and for each call to the merge function.

We have also used the *tareador_start_task* and *tareador_end_task* functions. The code is the following:

```

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        taredor_start_task("mu1");
        multisort(n/4L, &data[0], &tmp[0]);
        taredor_end_task("mu1");

        taredor_start_task("mu2");
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        taredor_end_task("mu2");

        taredor_start_task("mu3");
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        taredor_end_task("mu3");

        taredor_start_task("mu4");
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        taredor_end_task("mu4");

        taredor_start_task("me1");
        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        taredor_end_task("me1");

        taredor_start_task("me2");
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);
        taredor_end_task("me2");

        taredor_start_task("me3");
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
        taredor_end_task("me3");

    } else {
        // Base case
        basicsort(n, data);
    }
}

```

Figure 2: Modified version of the multisort function.

2.3 Task dependency graph (TDG)

We can clearly see that Figure 3 is divided into two parts: the multisort and the merge.

As we said before, in the multisort function we divide the input vector into 4 smaller vectors, until the vector size is smaller than $MIN_SORT_SIZE \times 4L$. This can be seen at the top of the figure, as we can see 4 different boxes (green, red, yellow and purple). Moreover, there is no data sharing between them, so we can execute them at the same time.

Then, there is a new recursion level inside each box. These four calls reach the limit value, so they do a *basicmerge* call. Afterwards, we need the first two childs to terminate before executing the first merge call (and the same for the third and

fourth child with the second merge call). Thus, there exist a dependence between the recursive calls to multisort and the calls to the first two merge functions. Finally, the third merge call can be executed only when the previous two merge calls have terminated (another dependence).

On the other hand, in the merge function, we create tasks when the number of elements of the vector that we want to sort is bigger or equal than $MIN_MERGE_SIZE \times 2L$. We see in the figure that the two created tasks do not have any dependence between them. A *basicmerge* call will be done when we reach the limit value.




Figure 3: Task dependency graph obtained using *Tareador*.

With all these information, we can conclude that we can use a **Tree strategy** (generate a task for each recursive function) because there are no dependences between sibling tasks. Besides, we can also use a **Leaf strategy** (generate a task in each base case of the functions).

Furthermore, we must be careful about the synchronization we need between the multisort tasks and the merge ones because, as we said before, the first merge call depends on the first two multisort calls, the second merge call depends on the third and fourth multisort calls and the third merge call depends on the previous two merge calls. Without synchronization we would have wrong results due to data racing.

2.4 Parallel Performance and Scalability

The following table shows the execution time and speed-up values obtained when varying the number of processors.

Processors	Execution Time [ns]	Speed-Up
1	20334411001	1
2	10173716001	1.99872013323365
4	5086725001	3.99754478510288
8	2550595001	7.97241858979085
16	1289922001	15.7640624667507
32	1289909001	15.7642213406029
64	1289909001	15.7642213406029

Table 1: Execution Time and Speed-Up varying the number of processors.

We can see that until 16 processors, the results obtained are pretty close to the ideal ones. However, from 16 to infinite processors there is no improvement on the speed-up values. This is because in Figure 3 we see that the maximum number of tasks executed at the same time are 16. For this reason, using more than 16 processors will not have any impact on the execution time and speed-up. See the traces extracted from *Paraver* in Annex 7.1.




Figure 4: Execution Time and Speed-Up plots varying the number of processors.

3 Parallelization and performance analysis with tasks

As we said in the previous section, we concluded that we can use both Tree and Leaf strategies for the parallelization of the multisort and merge functions.

In this section we are going to analyse the parallelization and performance of the 2 strategies implemented using the *OpenMP* clauses. The different versions of the code can be found in the next subsections. Nevertheless, the following code will be the same for the implementation of both strategies:

```

#pragma omp parallel
#pragma omp single
multisort(N, data, tmp);

```

Figure 5: Modified fragment of the given code in the main function.

3.1 Leaf Strategy

In this strategy, we define a task for the invocations of `basicsort` and `basicmerge` functions. That is, we create a task each time we are in the base case of `multisort` and `merge` functions.

In order to do it, we have used the `#pragma omp task` clause. As we have seen in the section 2.3, there exists dependencies between some function calls, so we will need an *OpenMP* clause to synchronize these functions (`#pragma omp taskwait`).

The code can be found in the `multisort-omp-leaf.c` file inside the codes directory. However, the important changes of the code are shown below.

```

void merge(long n, T left[n], T right[n], T result[n*2], long start,
           long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        #pragma omp task
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        merge(n, left, right, result, start, length/2);
        merge(n, left, right, result, start + length/2, length/2);
    }
}

```

Figure 6: Merge function implementing the Leaf strategy.


```

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        multisort(n/4L, &data[0], &tmp[0]);
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);

        #pragma omp taskwait

        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);

        #pragma omp taskwait

        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    } else {
        // Base case
        #pragma omp task
        basicsort(n, data);
    }
}

```

Figure 7: Multisort function implementing the Leaf strategy.

We can see in Figure 8 that only thread 0 is the one that creates tasks (in the base case of the functions), while the other tasks are executing them. Moreover, we can see that thread 0 is also executing tasks. We think that this happens because it does not create tasks very usually, so it has time to execute some of the tasks inside the pool.

On the other hand, we can also see the effect of the `#pragma omp taskwait` clause if we look at the first trace of the figure. There are plenty of red parts (synchronization parts) in all threads.




Figure 8: Fragment of the execution flow of the code using the Leaf strategy.

As in the Leaf strategy only one thread executes the recursive part of the functions, there is a big amount of time where the other threads are doing nothing. Consequently, the speed-up plots for the complete application and only the multisort function cannot be good.




Figure 9: Speed-up plot of the complete application using the Leaf strategy.




Figure 10: Speed-up plot of the multisort function using the Leaf strategy.

If we continue adding more threads to execute the code, this will result into an irrelevant improvement of both speed-up plots because there will be more threads that are doing no work, which means that the total amount of time in the Idle state will increase.

3.2 Tree Strategy

In this other strategy, we define a task during the recursive decomposition, when invoking multisort and merge recursive calls.

To do it, we have also used the `#pragma omp task` clause. However, we cannot use `#pragma omp taskwait` in the multisort function anymore because we have to stop the flow of the code until a group of recursive calls which some are siblings and others are not terminate. To achieve this, we have used the `#pragma omp taskgroup` clause. We have done 3 groups taking into account the dependences in Figure 3: multisort calls, first two merge calls and the third merge call.

The following code can be found in the `multisort-omp-tree.c` file inside the codes directory.

```

void merge(long n, T left[n], T right[n], T result[n*2], long start,
          long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp task
        merge(n, left, right, result, start, length/2);

        #pragma omp task
        merge(n, left, right, result, start + length/2, length/2);
    }
}

```

Figure 11: Merge function implementing the Tree strategy.

```

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        #pragma omp taskgroup
        {
            #pragma omp task
            multisort(n/4L, &data[0], &tmp[0]);

            #pragma omp task
            multisort(n/4L, &data[n/4L], &tmp[n/4L]);

            #pragma omp task
            multisort(n/4L, &data[n/2L], &tmp[n/2L]);

            #pragma omp task
            multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);
        }

        #pragma omp taskgroup
        {
            #pragma omp task
            merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);

            #pragma omp task
            merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L],
                0, n/2L);
        }

        #pragma omp task
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);
    } else {
        // Base case
        basicsort(n, data);
    }
}

```

Figure 12: Multisort function implementing the Tree strategy.

Now, with this new strategy, all threads can execute and create tasks. At the beginning only thread 0 could create tasks, but each time it creates one, the number of threads that can create tasks increases. As a consequence, the execution time will be reduced a lot because now the total amount of time in the Idle state is reduced.

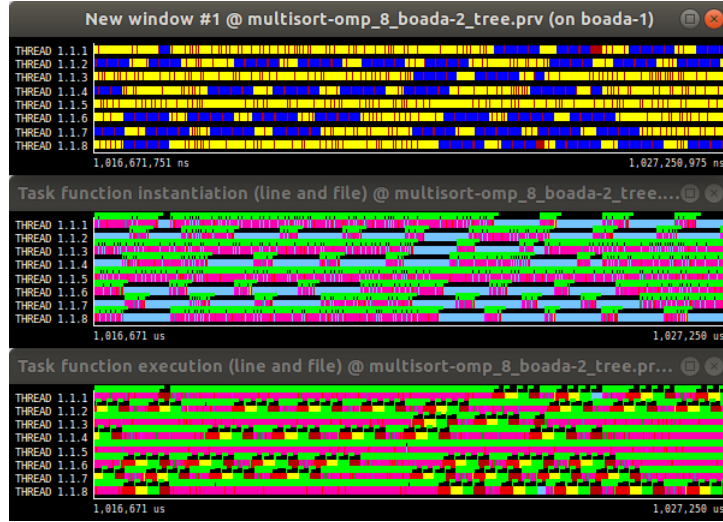


Figure 13: Fragment of the execution flow of the code using the Tree strategy.

As we have said before, the execution time is reduced a lot, improving the speed-up plots of both complete application and multisort function only. Hence, both plots will be much better than the ones obtained using the Leaf strategy.

Besides, we can see that this improvement has had a much impact on the multisort function only plot than in the complete application one.




Figure 14: Speed-up plot of the complete application using the Tree strategy.




Figure 15: Speed-up plot of the multisort function using the Tree strategy.

If we continue adding more threads to execute the code, this will result into a better performance of the code. We would have better execution time and speed-up results.

3.3 Task Cut-Off Mechanism

As we have seen, in the Tree strategy there is not a maximum recursion level, so if we have a large vector input we can be calling functions recursively for a long time.

To solve this problem, we can set a maximum recursion level using a cut-off mechanism that controls it for task generation (and their granularity).

To do this, we have used the *final(condition)* clause. When the condition is true, the task has the attribute final to true. The *omp_in_final()* function uses this attribute to decide whether it can create more tasks or not.

This code can be found in the *multisort-omp-tree-cutoff.c* file inside the codes directory.

```
int CUTOFF;

void merge(long n, T left[n], T right[n], T result[n*2], long start,
           long length, int depth) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        if (!omp_in_final()) {
            #pragma omp task final(depth >= CUTOFF)
            merge(n, left, right, result, start, length/2, depth + 1);

            #pragma omp task final(depth >= CUTOFF)
            merge(n, left, right, result, start + length/2, length/2,
                  depth + 1);
        } else {
            merge(n, left, right, result, start, length/2, depth + 1);
            merge(n, left, right, result, start + length/2, length/2,
                  depth + 1);
        }
    }
}
```

Figure 16: Merge function implementing the Tree strategy using the cut-off mechanism.

```

void multisort(long n, T data[n], T tmp[n], int depth) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition

        if (!omp_in_final()) {
            #pragma omp taskgroup
            {
                #pragma omp task final(depth >= CUTOFF)
                multisort(n/4L, &data[0], &tmp[0], depth + 1);

                #pragma omp task final(depth >= CUTOFF)
                multisort(n/4L, &data[n/4L], &tmp[n/4L], depth + 1);

                #pragma omp task final(depth >= CUTOFF)
                multisort(n/4L, &data[n/2L], &tmp[n/2L], depth + 1);

                #pragma omp task final(depth >= CUTOFF)
                multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L],
                    depth + 1);
            }

            #pragma omp taskgroup
            {
                #pragma omp task final(depth >= CUTOFF)
                merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L,
                    depth + 1);

                #pragma omp task final(depth >= CUTOFF)
                merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L],
                    0, n/2L, depth + 1);
            }

            #pragma omp task final(depth >= CUTOFF)
            merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n,
                depth + 1);
        } else {
            multisort(n/4L, &data[0], &tmp[0], depth + 1);
            multisort(n/4L, &data[n/4L], &tmp[n/4L], depth + 1);
            multisort(n/4L, &data[n/2L], &tmp[n/2L], depth + 1);
            multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L], depth + 1);

            merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L,
                depth + 1);
            merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0,
                n/2L, depth + 1);

            merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n,
                depth + 1);
        }
    } else {
        // Base case
        basicsort(n, data);
    }
}

```

Figure 17: Multisort function implementing the Tree strategy using the cut-off mechanism.

The CUTOFF variable can be set using the **-c cutoff** argument in the command line. If we do this, CUTOFF will be the maximum recursion level permitted.

To understand the cut-off mechanism, we have used a CUTOFF value of 0 and the *Paraver* tool to see the execution flow. The results are shown below.



Figure 18: Fragment of the execution flow of the code using the Tree strategy and CUTOFF = 0.

As the maximum recursion level is 0, we only create one task for each call function inside multisort. These new tasks will not be able to create tasks, and will execute the code in sequential mode.

In the recursive part, the multisort function does 7 calls (4 for the multisort and 3 for the merge). We can see this 7 created tasks in the third trace in Figure 18. The four large executions correspond to the multisort functions. Once the first two terminate, the first merge call (yellow one) can be executed. The same happens with the second merge call (red one) when the other two multisort functions end. Finally, the third merge function can start when the other two merge calls have terminated. This happens because of the taskgroup clause.

Once we have understood the cut-off mechanism, we are going to explore the performance of the code when using different values for the maximum recursion level. To do this, we have used the **submit-cutoff-omp.sh** script.




Figure 19: Execution time of the code when varying the recursion size.

With a recursion size value of 3 we get the best performance of the program. Analysing the scalability of the program using this optimum value, we see that the speed-up of the multisort function only is the best in comparison with the other two speed-up plots (Leaf and Tree strategies). However, the difference with the plot of the Tree strategy without the cut-off mechanism is very small.




Figure 20: Speed-up plot of the complete application using the Tree strategy and the cut-off mechanism.




Figure 21: Speed-up plot of the multisort function using the Tree strategy and the cut-off mechanism.

4 Parallelization and performance analysis with dependent tasks

In the previous sections we used `#pragma omp taskwait` and `#pragma omp taskgroup` to synchronize the tasks. However, it may happen that the execution time of one task is much greater than the other ones. As a result, we will lose time waiting that task to terminate.

A solution to this problem is to get rid of these two *OpenMP* clauses and use the *depend* clause to express dependencies. With this new clause, tasks do not have to wait until all previous tasks end but only the tasks that create their dependencies.

The new code is in *multisort-omp-tree-depend.c* file inside the codes directory.

```
void merge(long n, T left[n], T right[n], T result[n*2], long start,
           long length) {
    if (length < MIN_MERGE_SIZE*2L) {
        // Base case
        basicmerge(n, left, right, result, start, length);
    } else {
        // Recursive decomposition
        #pragma omp task
        merge(n, left, right, result, start, length/2);

        #pragma omp task
        merge(n, left, right, result, start + length/2, length/2);

        #pragma omp taskwait
    }
}
```

Figure 22: Merge function implementing the Tree strategy using the dependence clauses.

```

void multisort(long n, T data[n], T tmp[n]) {
    if (n >= MIN_SORT_SIZE*4L) {
        // Recursive decomposition
        #pragma omp task depend(out:data[0])
        multisort(n/4L, &data[0], &tmp[0]);

        #pragma omp task depend(out:data[n/4L])
        multisort(n/4L, &data[n/4L], &tmp[n/4L]);

        #pragma omp task depend(out:data[n/2L])
        multisort(n/4L, &data[n/2L], &tmp[n/2L]);

        #pragma omp task depend(out:data[3L*n/4L])
        multisort(n/4L, &data[3L*n/4L], &tmp[3L*n/4L]);

        #pragma omp task depend(in:data[0], data[n/4L])
                        depend(out:tmp[0])
        merge(n/4L, &data[0], &data[n/4L], &tmp[0], 0, n/2L);

        #pragma omp task depend(in:data[n/2L], data[3L*n/4L])
                        depend(out:tmp[n/2L])
        merge(n/4L, &data[n/2L], &data[3L*n/4L], &tmp[n/2L], 0, n/2L);

        #pragma omp task depend(in:tmp[0], tmp[n/2L])
        merge(n/2L, &tmp[0], &tmp[n/2L], &data[0], 0, n);

        #pragma omp taskwait
    } else {
        // Base case
        basicsort(n, data);
    }
}

```

Figure 23: Multisort function implementing the Tree strategy using the dependence clauses.

The first merge call cannot be executed until the previous tasks with *depend(out: data[0])* and *depend(out: data[n/4L])* terminate. Now, even though the third or fourth multisort calls have not terminated, the first merge call can be executed. The same happens for the rest of the calls with the depend clause. Thus, we will reduce a bit the execution time of the program.

Nevertheless, we must use *#pragma omp taskwait* at the end of both multisort and merge functions. Without them, the sorting would not be correct because it could happen that while the last merge call inside the multisort function is being executed, we terminate this function and start with other tasks. The next task would think that the vector is already sorted although that merge call would have not terminated yet. The same happens inside the merge function.

The most important thing we see in Figure 24 is that there is even less parts in Idle time than in the Tree flow (Figure 13). Hence, the execution time is reduced a bit with respect to the one using the Tree strategy.

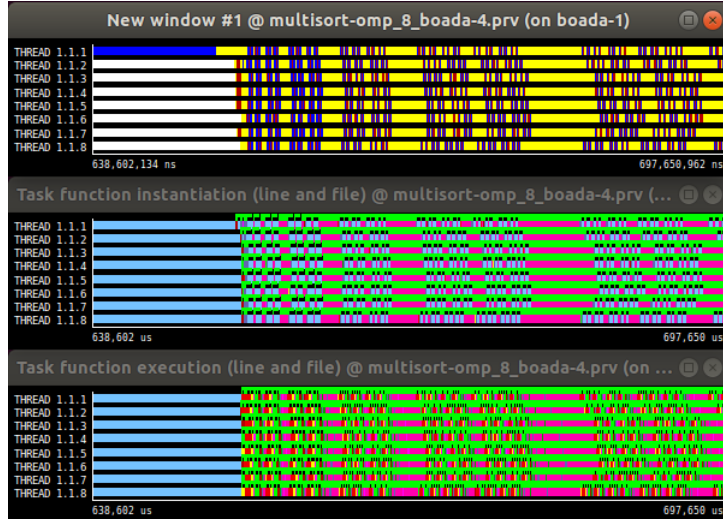


Figure 24: Fragment of the execution flow of the code using the Tree strategy and the depend clause.

Finally, in the scalability plots of this program we can observe that both have the best speed-ups with respect to all previous speed-up plots. This has been possible with the use of the depend clause.




Figure 25: Speed-up plot of the complete application using the Tree strategy and the cut-off mechanism.




Figure 26: Speed-up plot of the multisort function using the Tree strategy and the cut-off mechanism.

5 Optionals

5.1 Optional 1

In this part we used the code from the section Tree Strategy, to analyse the scalability plots on the other node types in boada. To do that, we edited for each node type

the maximum number of cores to be used in the **submit-strong-omp.sh** script (variable `np_NMAX`).

5.1.1 Boada-1 to 4

As we have seen in the first laboratory, the number of cores of these nodes is 12. Executing the script in one of those nodes we obtained the following Speed-up plots, which are the same than in the Tree Strategy section.




Figure 27: Speed-up plot of the complete application using the Tree strategy in Boada 3.




Figure 28: Speed-up plot of the multisort function using the Tree strategy in Boada 3.

5.1.2 Boada-5

In this node we have a total of 12 cores, which is the same that in Boada-1 to 4. However, in this node we have a different type of architecture. Submitting the script we obtained the following results:




Figure 29: Speed-up plot of the complete application using the Tree strategy in Boada 5.




Figure 30: Speed-up plot of the multisort function using the Tree strategy in Boada 5.

We can see that this architecture is much better than the one in Boada-1 to 4 because there has been an improvement on both Speed-up plots. Besides, the results in the second plot are almost the ideal case (discontinuous line).

5.1.3 Boada-6 to 8

These other nodes have a total of 16 cores. Thus, we have changed the value of np_NMAX to 16 in the **submit-strong-omp.sh** script. Afterwards we executed the script in one of these nodes.




Figure 31: Speed-up plot of the complete application using the Tree strategy in Boada 7.




Figure 32: Speed-up plot of the multisort function using the Tree strategy in Boada 7.

Although this node type is the one that has more cores, we do not get the best

results. Hence, this architecture is worse than the one in Boada-5. Nevertheless, we can see a little improvement in both plots with respect to the plots in Boada-1 to 4.

To sum up, we have seen that the best results are in the Boada-5 due to the architecture of this node, which has the best maximum core frequency.

5.2 Optional 2

In this other optional we had to parallelize the initialization of the data and temp vectors. The changes we made are shown below.

The following code can be found in the *multisort-omp-op2.c* file in the codes directory.

```
static void initialize(long length, T data[length]) {
    long i;
    #pragma omp for schedule(dynamic)
    for (i = 0; i < length; i++) {
        data[i] = rand();
    }
}

static void clear(long length, T data[length]) {
    long i;
    #pragma omp for schedule(dynamic)
    for (i = 0; i < length; i++) {
        data[i] = 0;
    }
}
```

Figure 33: Parallelization of Initialize and clear functions.

Before modifying the code, in the initialize function, each position of the vector data was computed using the value from the previous position of its vector. Thus, we could not parallelise it using the same code due to dependencies. In order to be able to parallelize the initialize function, we have changed the way the random numbers are created, now each number is completely random. We have used the **dynamic** schedule type in order to balance the work in all threads.

Afterwards we have submitted the scripts **submit-strong-omp.sh** and **submit-omp-i.sh**, obtaining the following speed-up plots and *Paraver* traces.




Figure 34: Speed-up plot of the complete application with the parallelisation of the initialize and clear function and using the Tree strategy.




Figure 35: Speed-up plot of the multisort function with the parallelisation of the initialize and clear function and using the Tree strategy.

We though that there would have been a big improvement in the speed-up plot of the complete application. However, it has not been a big improvement but a good one. Probably trying to parallelise other parts of the code would result in a bigger improvement on the speed-up plot of the complete application.

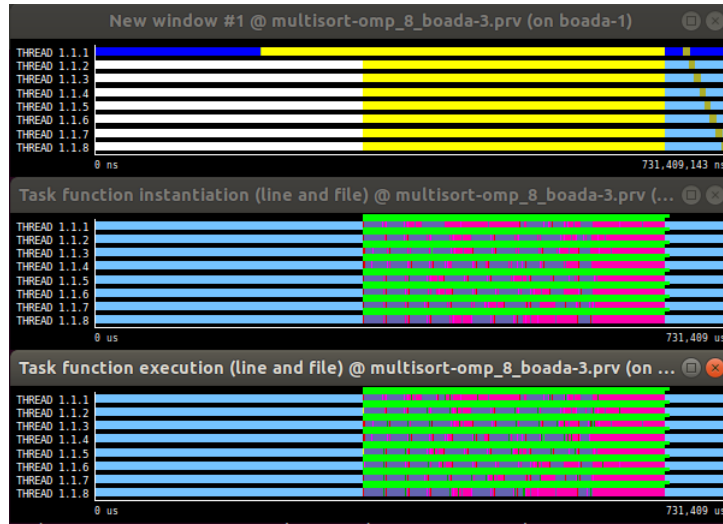


Figure 36: Execution flow of the modified code using 8 threads.

In the *Paraver* traces we can see that the parallel part of the program has increased a bit, which is what we expected.

To conclude, we can see that we have better results when the initializations are parallelized, especially in the complete application Speed-up plot. In the execution flow we have seen that the serial part has been reduced.

6 Conclusion

In this three laboratory sessions we have seen different strategies for the mergesort algorithm.

With *Tareador* we have seen that it is possible to implement both Leaf and Tree strategies due to the dependencies.

On the one hand, the Leaf strategy has not given use the best results. This is because it only creates a task for each leaf of the tree generated with the recursive calls. Hence, in almost all the time there is only one thread doing work.

On the other hand, the Tree strategy has given us the best results because all threads can create tasks. As a consequence, the total amount of time of threads in Idle state (doing no work) has decreased a lot. Thus, the performance of this strategy is better than the first one.

When using the cut-off mechanism, we saw that the execution time varies when changing the value of the recursion size. Using the optimum value we have observed that we improved the performance of the Tree strategy.

We must remark that using the depend clause we have got better results because tasks will not be suspended until all previous tasks finishes, they will be suspended until the variables they need have been correctly modified.

Finally, in the first optional we have seen that it is not important the number of cores to use but the type of architecture of the node. Even though Boada-5 has fewer cores than Boada-6 to 8, it has better results due to its architecture type.

In the second optional, we have seen that parallelising the initialization of the vectors has increased a bit the performance of the complete application. Nevertheless, we think that we would get better results if we parallelise other functions (e.g. the `check_sorted` function).

7 Annex

7.1 Scalability anaysis: Traces using different number of processors




Figure 37: Execution flow of the code using 1 processor.




Figure 38: Execution flow of the code using 2 processors.




Figure 39: Execution flow of the code using 4 processors.




Figure 40: Execution flow of the code using 8 processors.




Figure 41: Execution flow of the code using 16 processors.




Figure 42: Execution flow of the code using 32 processors.




Figure 43: Execution flow of the code using 64 processors.